

Observing Node Performance with DTrace

- **Introduction**
 - What is DTrace?
 - How DTrace can be used with node.js
- **DTrace kernel actions**
- **DTrace the node engine**
- **DTrace node applications**

What is DTrace?



- **Tool that allows one to dynamically instrument code from application level and into the kernel.**
- **Can be used safely on production systems.**
- **Uses:**
 - Performance Analysis
 - Debugging
 - Code coverage
 - Find out wtf is happening in your software
- **Available on illumos, smartOS, and other Solaris 10 derivatives, as well as *BSD and Mac OS X.**

- **System Call** - Request for an action by the Operating System
- **Probe** - An instrumentation point in the code
 - Dynamic and Static probes are provided, and new ones can be added
 - A probe is specified by a 4-tuple:
 - `provider:module:function:probename{action}`
- **Action** - Executed when a probe fires
- **Predicate** - Optional boolean to determine whether or not to execute the action
- **Example:** `syscall::read:entry/pid == 713/{trace();}`


Node.js with DTrace Support



- From www.nodejs.org download site

```
# curl -O http://nodejs.org/dist/v0.8.11/node-v0.8.11.tar.gz
% Total      % Received % Xferd  Average Speed   Time    Time     Time  Current
           % Dload    % Upload   Total   Spent    Left     Speed

100 11.2M  100 11.2M    0     0   253k      0  0:00:45  0:00:45 --:--:--  349k
# gtar -xpf node-v0.8.11.tar.gz
# pkgin install gcc-compiler-4.6.1
...
# cd node-v0.8.11
# ./configure
...
# make
...
# make install    <-installs in /usr/local/bin
...
# export PATH=/usr/local/bin:$PATH
# node -v
v0.8.11
# cd ..
# npm install dtrace-provider
...
# npm install restify    <- This is not necessary, but will be used in some of the demos
...
```



- **With DTrace, you can trace events in**
 - The node Engine
 - Node.js scripts
 - The kernel (system calls, scheduling, memory management, etc.)

- **Show system calls made by a running node process**

```
# dtrace -n 'syscall:::entry/pid==26442/{}'  
dtrace: description 'syscall:::entry' matched 234 probes  
CPU      ID      FUNCTION:NAME  
  1    10157      write:entry  
  1    10283    lwp_park:entry  
  1    10155      read:entry  
  4    10155      read:entry  
  4    10157      write:entry  
...
```

- **Count system calls made by a running node process**

```
# dtrace -n 'syscall:::entry/execname == "node"/{@[probefunc]=count();}'  
dtrace: description 'syscall:::entry' matched 234 probes  
(^C)  
...  
munmap      31  
portfs      36  
lwp_park     37  
fcntl       39  
mmap64      53  
#
```

An Example Measuring System Call Latency



- **systeme.d**

```
#!/usr/sbin/dtrace -s

#pragma D option quiet

syscall::entry
/execname == "node"/
{
    self->ts = timestamp;
}

syscall::return
/self->ts/
{
    @[probfunc] = quantize(timestamp - self->ts);
    self->ts = 0;
}

END
{
    printa("SYSCALL      NSECS                      # OF OCCURANCES\n%s%@1x\n", @);
}
```



An Example Measuring System

Call Latency (Continued)



```
# ./system.d
...
SYSCALL      NSECS      # OF OCCURANCES
read
value  ----- Distribution ----- count
1024   |
2048   | @@@@@@@@@@@@@@@@
4096   | @@@@@@@@@@
8192   | @@@@@@@@@@
16384  | @@@@
32768  |
65536  |
131072 |
262144 |
524288 |
1048576 |
2097152 |
4194304 |
8388608 |
16777216 |
33554432 |
67108864 |
134217728 | @@@@
268435456 |
...

```



- **With DTrace, you can trace events in**
 - The node Engine
 - Node.js scripts
 - The kernel (system calls, scheduling, memory management, etc.)

The Node DTrace Provider



- Set of USDT probes built into node

```
# dtrace -l -n 'node*:::{'
```

ID	PROVIDER	MODULE	FUNCTION NAME
57166	node11665	node	
		_ZN4nodeL14dtrace_gc_doneEN2v86GCTypeENS0_15GCCallbackFlagsE	gc-done
57167	node11665	node	
		_ZN4nodeL15dtrace_gc_startEN2v86GCTypeENS0_15GCCallbackFlagsE	gc-start
57168	node11665	node	_ZN4node26DTRACE_HTTP_CLIENT_REQUESTERKN2v89ArgumentsE
			http-client-request
57169	node11665	node	_ZN4node27DTRACE_HTTP_CLIENT_RESPONSEERKN2v89ArgumentsE
			http-client-response
57170	node11665	node	_ZN4node26DTRACE_HTTP_SERVER_REQUESTERKN2v89ArgumentsE
			http-server-request
57171	node11665	node	_ZN4node27DTRACE_HTTP_SERVER_RESPONSEERKN2v89ArgumentsE
			http-server-response
57172	node11665	node	_ZN4node28DTRACE_NET_SERVER_CONNECTIONERKN2v89ArgumentsE
			net-server-connection
57173	node11665	node	_ZN4node22DTRACE_NET_SOCKET_READERKN2v89ArgumentsE
			net-socket-read
57174	node11665	node	_ZN4node23DTRACE_NET_SOCKET_WRITEERKN2v89ArgumentsE
			net-socket-write
57175	node11665	node	_ZN4node21DTRACE_NET_STREAM_ENDERKN2v89ArgumentsE
			net-stream-end

The Node DTrace Provider

Probe Arguments



```
# dtrace -l -v -n 'node*:::http-server-request, node*:::http-server-response{'
```

ID	PROVIDER	MODULE	FUNCTION NAME
57170	node11665	node _ZN4node26DTRACE_HTTP_SERVER_REQUESTERKN2v89ArgumentsE	http-server-request

Probe Description Attributes

Identifier Names: Private

Data Semantics: Private

Dependency Class: Unknown

Argument Attributes

Identifier Names: Evolving

Data Semantics: Evolving

Dependency Class: ISA

Argument Types

args[0]: node_http_request_t *

args[1]: node_connection_t *

57171	node11665	node _ZN4node27DTRACE_HTTP_SERVER_RESPONSEERKN2v89ArgumentsE	http-server-response
-------	-----------	--	-----------------------------

...

args[0]: node_connection_t *

The Node DTrace Provider

Probe Arguments (Continued)



- In `node-v0.8.11/src/node.d`

```
typedef struct {  
    string url;  
    string method;  
    string forwardedFor;  
} node_http_request_t;  
  
typedef struct {  
    int fd;  
    string remoteAddress;  
    int remotePort;  
    int bufferSize;  
} node_connection_t;
```

The Node DTrace Provider: Example 1



```
/* echo-server.d */

#pragma D option quiet

BEGIN
{
    printf("%-22s %-20s %-8s %-16s %-16s %-16s\n",
        "DIRECTION", "URL", "METHOD", "REMOTEADDRESS", "REMOTEPORT", "BUFFERSIZE");
}

node*::http-server-request
{
    printf("%-22s %-20s %-8s %-16s %-16d %-16d\n",
        probename, args[0]->url, args[0]->method, args[1]->remoteAddress,
        args[1]->remotePort, args[1]->bufferSize);
}

node*::http-server-response
{
    printf("%-22s %-20s %-8s %-16s %-16d %-16d\n",
        probename, " ", " ", args[0]->remoteAddress,
        args[0]->remotePort, args[0]->bufferSize);
}
```

The Node DTrace Provider: Example 1 (Continued)



- Client

```
# curl http://165.225.154.78:8080/echofile-server.js > /dev/null
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           %             Dload  Upload  Total  Spent    Left   Speed
100  1377  100  1377    0     0   382k      0  --:--:-- --:--:-- --:--:--  672k
```

- Server

```
# dtrace -L /usr/local/lib/dtrace -s echo-server.d
DIRECTION          URL                METHOD  REMOTEADDRESS      REMOTEPORT
BUFFERSIZE
http-server-request /echofile-server.js GET      62.203.55.164      58027          0
http-server-response      62.203.55.164      58027          0
http-server-response      62.203.55.164      58030          0
http-server-request /echofile-server.js GET      62.203.55.164      58030          0
http-server-request /echofile-server.js GET      62.203.55.164      58036          0
http-server-response      62.203.55.164      58036          0
http-server-request /echofile-server.js GET      62.203.55.164      58037          0
http-server-response      62.203.55.164      58037          0
http-server-request /echofile-server.js GET      62.203.55.164      58038          0
http-server-response      62.203.55.164      58038          0
http-server-request /systime.d         GET      62.203.55.164      58363          0
http-server-response      62.203.55.164      58363          0
http-server-request /favicon.ico       GET      62.203.55.164      58364          0
http-server-response      62.203.55.164      58364          0
...
```


Request/Response Latency



```
/* server-latency.d */

#pragma D option quiet

node*::http-server-request
{
    ts[args[1]->remoteAddress, args[1]->remotePort] = timestamp;
    url[ts[args[1]->remoteAddress, args[1]->remotePort]] = args[0]->url;
}

node*::http-server-response
/ts[args[0]->remoteAddress, args[0]->remotePort]/
{
    this->t = ts[args[0]->remoteAddress, args[0]->remotePort];
    @[url[this->t], args[0]->remoteAddress] = quantize((timestamp-this->t)/1000);
    ts[args[0]->remoteAddress, args[0]->remotePort] = 0;
}

END
{
    printf("%-20s: %-16s\n", "URL", "REMOTEADDRESS");
    printa("%-20s: %-16s\nMICROSECONDS\n%@", @);
}
```


Request/Response Latency (Continued)



```
# dtrace -L /usr/local/lib/dtrace -s server-latency.d
```

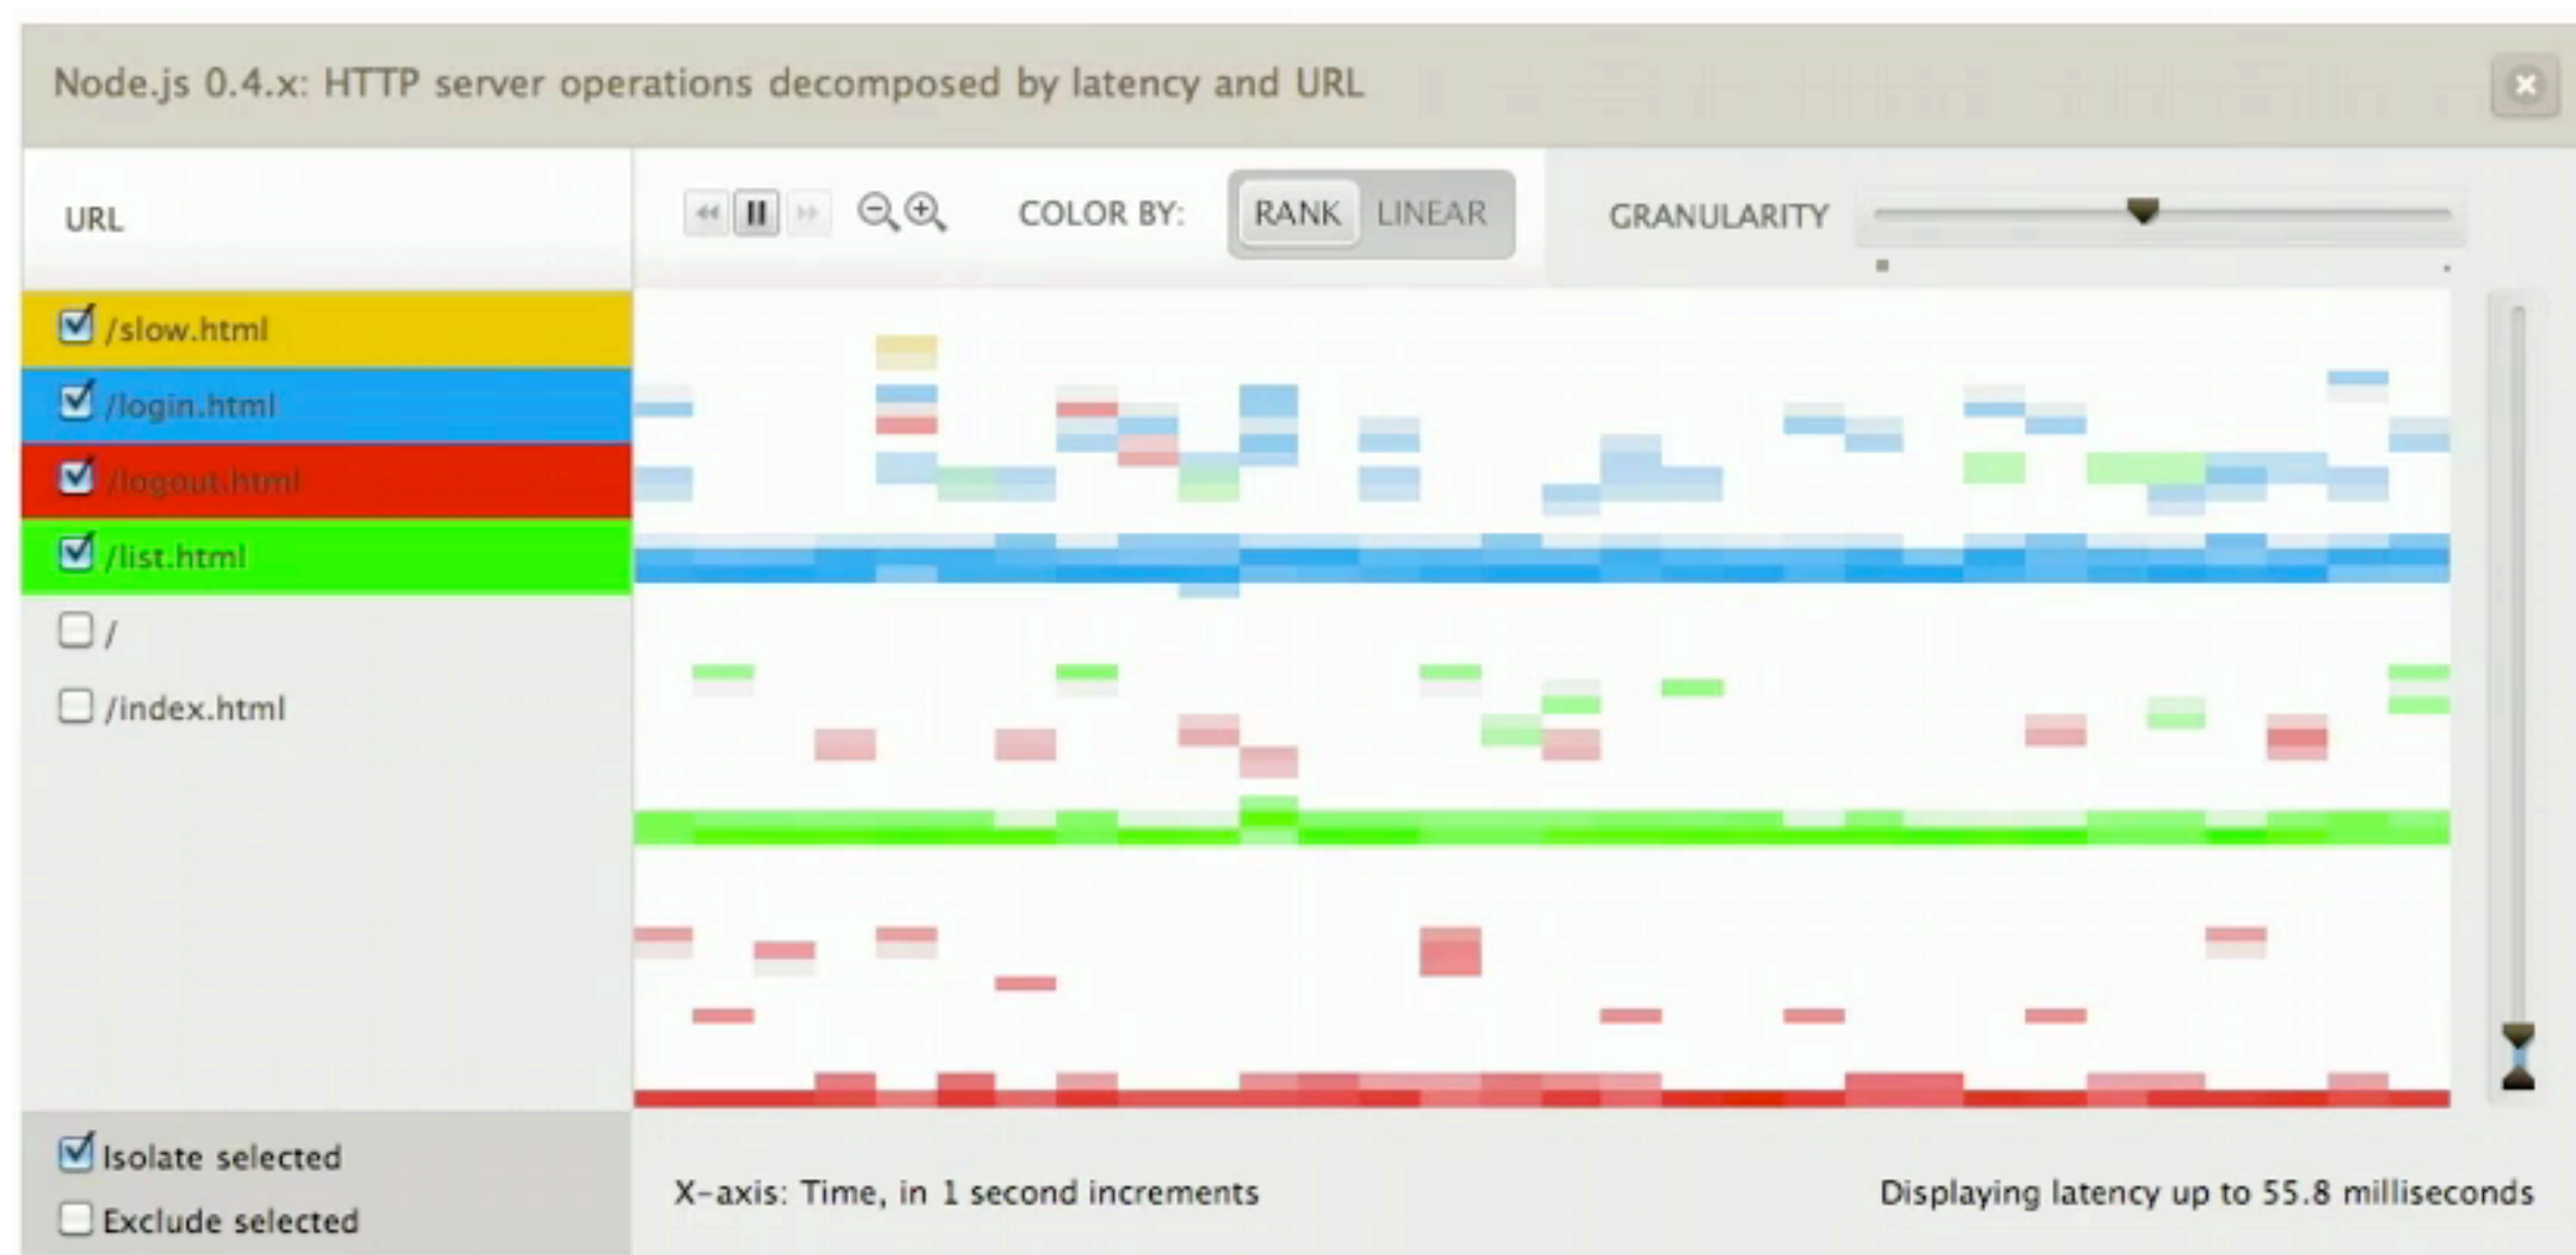
```
URL           : REMOTEADDRESS
/tmp/words    : 165.225.154.77
MICROSECONDS
```


value	----- Distribution -----	count
1024		0
2048	@@@@	11
4096	@@@@@@@@@@@@@@@@@@	43
8192	@@@@@@	14
16384	@@@@@@@@@@@@@@	31
32768		1
65536		0

```
/tmp/words    : 83.79.36.187
MICROSECONDS
```

value	----- Distribution -----	count
524288		0
1048576	@	3
2097152	@@	4
4194304	@@	4
8388608	@@@@	11
16777216	@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@	74
33554432	@@	4
67108864		0

Heatmaps





- **With DTrace, you can trace events in**
 - The node Engine
 - Node.js scripts
 - The kernel (system calls, scheduling, memory management, etc.)

DTrace Your node.js Application Joyent

- The dtrace-provider for Node.js allows you to create statically defined probes (USDT) in your application.
- Effectively, a way to add print statements to your scripts which only have effect when/if the probes are enabled.
- But better than print... You decide what to enable and what to print at runtime.
- Install
 - `npm install dtrace-provider`

Add Probes to Your Node App



```
/* echofile-server.js */
...
var dtp = require('dtrace-provider').createDTraceProvider('echofile-
server');

dtp.addProbe('echo-start', 'char *');
dtp.addProbe('echo-done', 'char *', 'int');
dtp.addProbe('echo-error', 'char *', 'char *')
...

dtp.fire('echo-start', function() {
    return [req.params[0]];
});
...
dtp.fire('echo-error', function() {
    return [req.params[0], JSON.stringify(e)];
});
...
dtp.fire('echo-done', function() {
    return [req.params[0], len];
});
...
```

- **Define probes and arguments**

- **Add probes to your code**

DTrace The Added Probes



```
#!/usr/sbin/dtrace -s
```

```
#pragma D option quiet
```

```
echofile-server*:::echo-start
```

```
{  
    printf("%s: %s\n", probename, copyinstr(arg0));  
}
```

```
echofile-server*:::echo-done
```

```
{  
    printf("%s: %s %d bytes\n", probename, copyinstr(arg0), arg1);  
}
```

```
echofile-server*:::echo-error
```

```
{  
    printf("%s\n", copyinstr(arg1));  
}
```

- **Use dtrace to enable the probes you've added**

Enabling the Added Probes



```
# ./echofile-server.d
echo-start: tmp/bigwords
echo-done: tmp/bigwords 20667400 bytes
echo-start: tmp
echo-done: tmp 116 bytes
{"errno":28,"code":"EISDIR"}
echo-start: blah
{"errno":34,"code":"ENOENT","path":"blah"}
...
```


List Probes Built-in for Restify



```
# dtrace -l -P 'myapp*'
ID PROVIDER MODULE FUNCTION NAME
57309 myapp13446 module func get100-start
57310 myapp13446 module func get100-done
57311 myapp13446 module func get100-
parseAccept-start
57312 myapp13446 module func get100-
parseAccept-done
57313 myapp13446 module func get100-
parseQueryString-start
57314 myapp13446 module func get100-
parseQueryString-done
57315 myapp13446 module func get100-parseBody-
start
57316 myapp13446 module func get100-parseBody-
done
57317 myapp13446 module func get100-sget-start
57318 myapp13446 module func get100-sget-done
```


- <https://github.com/mcavage/node-restify>
- <http://mcavage.github.com/presentations/dtrace-conf-2012-04-03/>
- <https://github.com/chrisa/node-dtrace-provider>
- <http://dtrace.org/blogs/blog/category/node-js/>
- <http://dtrace.org/blogs/dap/files/2012/05/fluent.pdf>
- <http://dtrace.org/blogs/bmc/2010/08/30/dtrace-node-js-and-the-robinson-projection/>
- <http://dtrace.org/blogs/dap/2012/01/05/where-does-your-node-program-spend-its-time/>
- <http://dtrace.org/blogs/brendan/2011/09/26/>

Acknowledgements



- Thanks to NodeDublin, Joyent Engineering (Bryan Cantrill, Mark Cavage, Robert Mustacchi, Dave Pacheco, and others), Marco Meinardi of Joyent
- Slides, node.js scripts, and D scripts are on <https://github.com/max123/NodeDublin-DTrace-talk.git>
- Thanks for listening!
- max@joyent.com, @mrbruning, mbruning.blogspot.com

